
2016/11/16, Wensday, Cloudy

by Jincheng Su @ Hikvision, email: jcsu14@fudan.edu.cn

Problem left yesterday

Pro1. Caffe installation not succeeded. It may be the problem of not installing cuDNN successfully.

Pro2. ``/path/to/bcnn/run_experiments_bcnn_train.m`` still won't work.'

Going to put aside the first problem for some days, solve the pro2 firstly with full efforts.

Keep on analysing the source code

Overview of `run_experiments_bcnn_train.m` :

```
Entry :`run_experiments_bcnn_train()`  
|  
Set Pretrained CNN(s): `bcnnmm`, `bcnnvdm`, or `bcnnvdvd` or all three  
|  
Set Model Parameters and dataset: `[opts, imdb] = model_setup(/*parameter pairs*/)`  
|  
Begin to train a bcnn model: `imdb_bcnn_train_dag(imdb, opts)`  
|  
End
```

Overview of `model_setup(...)` :

```
Entry: `[opts, imdb] = model_setup(varargin)`  
|  
Setup data structure `opts`: `batchSize`, `numEpochs`, `learningRate`, ...  
|  
Setup data structure `opts.encoders`: what the hell is encoders?  
|  
Load dataset: `cub`, `cars`, or `aircraft`, or others  
|  
End: save imdb
```

Overview of `imdb_bcnn_train_dag(imdb, opts)`

```
Entry: `imdb_bcnn_train_dag(imdb, opts)`  
|  
Setup some training params: `opts.train....`  
|  
Build up and Initialize network:  
`initializeNetworkSharedWeights` or `initializeNetworkTwoStreams`  
|  
Applying SGD algorithm to train the bcnn  
|  
End: save networks
```

How to build up the whole bcnn model?

```
Entry:
`net = initializeNetworkSharedWeights(imdb, encoderOpts, opts)`
|
Load the pretrained model:
`net = load(encoderOpts.modela)`
|
Truncate the Pretrained model for feature extraction:
`net.layers = net.layers(1:maxLayer)`
|
Add batch normalization
|
Stack bilinearpool layer:
`net.layers{end+1} = struct('type', 'bilinearpool', 'name', 'blp')`
|
Stack sqrt normalization:
`net.layers{end+1} = struct('type', 'sqrt', 'name', 'sqrt_norm')`
|
Stack l2 normalization:
`net.layers{end+1} = struct('type', 'l2norm', 'name', 'l2_norm')`
|
Stack a linear classifier netc:
`net.layers{end+1} = struct('type', 'conv', 'name', 'classifier', ...)`
|
Finally, add a softmaxloss layer:
`net.layers{end+1} = struct('type', 'softmaxloss', 'name', 'loss')`
|
Pretrain the linear classifier with logistic regression
```

So, how to set the pretrained model?

```
% run_experiments_bcnn_train.m

bcnnmm.name = 'bcnnmm' ;
bcnnmm.opts = {...
'type', 'bcnn', ...
'modela', 'data/models/imagenet-vgg-m.mat', ... % initialize network A with pre-trained model
'layera', 14,... % specify the output of certain layer of network A
to be bilinearly combined
'modelb', 'data/models/imagenet-vgg-m.mat', ... % initialize network B with pre-trained model
'layerb', 14,... % specify the output of certain layer of network B
to be bilinearly combined
'shareWeight', true,... % true: symmetric implementation where two network
s are identical
} ;

bcnnvdm.name = 'bcnnvdm' ;
bcnnvdm.opts = {...
'type', 'bcnn', ...
'modela', 'data/models/imagenet-vgg-verydeep-16.mat', ...
'layera', 30,...
'modelb', 'data/models/imagenet-vgg-m.mat', ...
'layerb', 14,...
'shareWeight', false,... % false: asymmetric implementation where two netwo
rks are distinct
} ;

bcnnvdvd.name = 'bcnnvdvd' ;
```

```

bcnnvdvd.opts = {...
'type', 'bcnn', ...
'modela', 'data/models/imagenet-vgg-verydeep-16.mat', ...
'layera', 30,...
'modelb', 'data/models/imagenet-vgg-verydeep-16.mat', ...
'layerb', 30,...
'shareWeight', true,...
};

setupNameList = {'bcnnmm'};
encoderList = {{bcnnmm}};
% setupNameList = {'bcnnvdvd'};
% encoderList = {{bcnnvdvd}};
datasetList = {'cub', 1}

```

Now that I have known that in order to setup a whole bcnn model, all that we need to do is to download the pretrained models `imagenet-vgg-m.mat` and `imagenet-vgg-verydeep-16.mat`. And then put the two models under `/path/to/bcnn_root/data/models` or change the path in `bcnnxxx.opts` to point to wherever your models are. Then setup variables `setupNameList = {'xxx'}` and `encoderList = {{xxx}}`, and `datasetList` as well. Done! That is seem to be all we need to do for building up a bcnn model.

Now that I am of quite certainty that the bcnn model should have been setup. The remained questions are

1. how to setup the dataset?
2. how to setup the training and testing params?

So how the dataset can be set up?

To find out the dataset setting up function, look into the function `[opts, imdb] = model_setup(varargin)` :

```

switch opts.dataset
case 'cubcrop'
    imdb = cub_get_database(opts.cubDir, true, false);
case 'cub'
    imdb = cub_get_database(opts.cubDir, false, opts.useVal);
case 'aircraft-variant'
    imdb = aircraft_get_database(opts.aircraftDir, 'variant');
case 'cars'
    imdb = cars_get_database(opts.carsDir, false, opts.useVal);
case 'imagenet'
    imdb = cnn_imagenet_setup_data('dataDir', opts.ilsvrDir);
case 'imagenet-224'
    imdb = cnn_imagenet_setup_data('dataDir', opts.ilsvrDir_224);
otherwise
    error('Unknown dataset %s', opts.dataset) ;
end

```

To use the birds dataset `cub`, look into its function `imdb = cub_get_database(cubDir, useCropped, useVal)` :

```

% cub_get_database.m

% the directory where the real images reside
if useCropped
    imdb.imageDir = fullfile(cubDir, 'images_cropped') ;
else
    imdb.imageDir = fullfile(cubDir, 'images');
end

```

```

...

% read the class names, image names, bounding boxes, labels...

[~, classNames] = textread(fullfile(cubDir, 'classes.txt'), '%d %s');
imdb.classes.name = horzcat(classNames(:));

% Image names
[~, imageNames] = textread(fullfile(cubDir, 'images.txt'), '%d %s');
imdb.images.name = imageNames;
imdb.images.id = (1:numel(imdb.images.name));

...

% if use validation, set 1/3 to validation set
if useVal
    rng(0)
    trainSize = numel(find(imageSet==1));

    trainIdx = find(imageSet==1);

    % set 1/3 of train set to validation
    valIdx = trainIdx(randperm(trainSize, round(trainSize/3)));
    imdb.images.set(valIdx) = 2;
end

...

```

Download the CUB-200-2011 dataset and unzip and find that all the required files are contained in the package, all we need to do is put the directory under `/path/to/bcnn_root/data` with a new name `cub`.

- Dataset details:
 - Birds: [CUB-200-2011 dataset](#). Birds + box uses bounding-boxes at training and test time.
 - Aircrafts: [FGVC aircraft dataset](#)
 - Cars: [Stanford cars dataset](#)
- These results are with domain specific fine-tuning. For more details see the updated [B-CNN tech report](#).

With `bcnn` model and `cub` dataset set up, all that remains is to check the related params to train a `bcnn` for CUB-200-2011 dataset.

Have a look at how the `model_setup(...)` is called:

```

% run_experiments_bcnn_train.m
[opts, imdb] = model_setup('dataset', dataset, ...
    'encoders', encoderList{ee}, ...
    'prefix', 'checkgpu', ... % output folder name
    'batchSize', 64, ...
    'imgScale', 2, ... % specify the scale of input images
    'bcnnLRinit', true, ... % do logistic regression to initialize softmax layer
    'dataAugmentation', {'f2', 'none', 'none'}, ... % do data augmentation [train, val, test]. Only support flipping for train set on current release.
    'useGpu', [1], ... %specify the GPU to use. 0 for using CPU
    'learningRate', 0.001, ...
    'numEpochs', 100, ...
    'momentum', 0.9, ...
    'keepAspect', true, ...
    'printDatasetInfo', true, ...
    'fromScratch', false, ...
    'rgbJitter', false, ...
    'useVal', false, ...
    'numSubBatches', 1);

```

It seems that all the related params can be set inside this function call. And seems we need to do nothing with it other than change `useVal` to `true`.

Anyway, it do not harm to have a look into the `model_setup` function.

```
% model_setup.m

function [opts, imdb] = model_setup(varargin)

% Copyright (C) 2015 Tsung-Yu Lin, Aruni RoyChowdhury, Subhransu Maji.
% All rights reserved.
%
% This file is part of the BCNN and is made available under
% the terms of the BSD license (see the COPYING file).

setup ;

opts.seed = 1 ;
opts.batchSize = 128 ;
opts.numEpochs = 100;
opts.momentum = 0.9;
opts.learningRate = 0.001;
opts.numSubBatches = 1;
opts.keepAspect = true;
opts.useVal = false;
opts.fromScratch = false;
opts.useGpu = 1 ;
opts.regionBorder = 0.05 ;
opts.numDCNNWords = 64 ;
opts.numDSIFTWords = 256 ;
opts.numSamplesPerWord = 1000 ;
opts.printDatasetInfo = false ;
opts.excludeDifficult = true ;
opts.datasetSize = inf;
% opts.encoders = {struct('name', 'rcnn', 'opts', {})} ;
opts.encoders = {} ;
opts.dataset = 'cub' ;
opts.carsDir = 'data/cars';
opts.cubDir = 'data/cub';
opts.aircraftDir = 'data/fgvc-aircraft-2013b';
opts.ilsvrcDir = '/home/tsungyulin/dataset/ILSVRC2014/CLS-LOC/';
opts.ilsvrcDir_224 = '/home/tsungyulin/dataset/ILSVRC2014/CLS-LOC-224/';
opts.suffix = 'baseline' ;
opts.prefix = 'v1' ;
opts.model = 'imagenet-vgg-m.mat';
opts.modela = 'imagenet-vgg-m.mat';
%opts.cropSize = 227/256; % added by jcsu

%opts.model = 'imagenet-vgg-verydeep-16.mat';
%opts.modela = 'imagenet-vgg-verydeep-16.mat';
opts.modelb = [];
opts.layer = 14;
opts.layera = [];
opts.layerb = [];
opts.imgScale = 1;
opts.bcnnLRinit = false;
opts.bcnnLayer = 14;
opts.rgbJitter = false;
```

```

opts.dataAugmentation = {'none', 'none', 'none'};
opts.cudnn = true;
opts.nonftbcnnDir = 'nonftbcnn';
opts.batchNormalization = false;
opts.cudnnWorkspaceLimit = 1024*1024*1204;

[opts, varargin] = vl_argparse(opts,varargin) ;

opts.expDir = sprintf('data/%s/%s-seed-%02d', opts.prefix, opts.dataset, opts.seed) ;
opts.nonftbcnnDir = fullfile(opts.expDir, opts.nonftbcnnDir);
opts.imdbDir = fullfile(opts.expDir, 'imdb') ;
opts.resultPath = fullfile(opts.expDir, sprintf('result-%s.mat', opts.suffix)) ;

opts = vl_argparse(opts,varargin) ;

if nargin <= 1, return ; end

...

```

Owh, Jesus! what the hell are all those params in `opts` !!

With a little thoughts, it seems that these are default values to the params. And we can change it via `argument pairs` in the `model_setup('name', 'value', ...)` call.

Anyway, just leave it alone.

In summary, what I have done are:

download `imagenet-vgg-m.mat` and `imagenet-vgg-verydeep-16.mat`, and put it under

`/path/to/bcnn_root/data/models` download datasets `CUB-200-2011`, and put it under `/path/to/bcnn_root/data` with a new name `cub`

*check the variables `setupNameList`, `encoderList`, and `datasetList` in `run_experiments_bcnn_train.m`, to use the `bcnnmm` model and `cub` dataset.

OK, run `run_experiments_bcnn_train` !

```

196: 196.House_Wren (train: 29, test: 30 total: 59)
197: 197.Marsh_Wren (train: 30, test: 30 total: 60)
198: 198.Rock_Wren (train: 30, test: 30 total: 60)
199: 199.Winter_Wren (train: 30, test: 30 total: 60)
200: 200.Common_Yellowthroat (train: 30, test: 30 total: 60)
Inf: **total** (train: 11788, test: 11788 total: 11788)
dataset: there are 11788 images (5994 trainval 5794 test)
collecting image stats: batch starting with image 2 ...Error using vl_argparse (line 101)
Unknown parameter 'cropSize'

Error in vl_argparse (line 80)
    opts = vl_argparse(opts, vertcat(params,values)) ;

Error in imdb_get_batch_bcnn (line 35)
    opts(i) = vl_argparse(opts(i), {varargin{1}(i),varargin{2:end}});

Error in getBatchSimpleNNWrapper>getBatchSimpleNN (line 10)
im = imdb_get_batch_bcnn(images, opts, ...

Error in getBatchSimpleNNWrapper>@(imdb,batch)getBatchSimpleNN(imdb,batch,opts) (line 4)
fn = @(imdb, batch) getBatchSimpleNN(imdb, batch, opts) ;

Error in imdb_bcnn_train_dag>getImageStats (line 213)
    temp = fn(imdb, batch) ;

```

```
Error in imdb_bcnn_train_dag (line 83)
    [averageImage, rgbMean, rgbCovariance] = getImageStats(imdb, bopts) ;

Error in run_experiments_bcnn_train (line 81)
    imdb_bcnn_train_dag(imdb, opts);
```

ERROR! ...

Well, set some breakpoints, and locate the error (NOTE: although MATLAB gives lots of helpful error information to help us debugging, chances are that the error location can be not that accurate. And it also helps a lot to trace back the execution stack.):

```
**function run_experiments_bcnn_train():**
    -> imdb_bcnn_train_dag(imdb, opts);
**function imdb_bcnn_train_dag(imdb, opts, varargin):**
    -> [averageImage, rgbMean, rgbCovariance] = getImageStats(imdb, bopts);
**function [averageImage, rgbMean, rgbCovariance] = getImageStats(imdb, opts):**
    -> fn = getBatchSimpleNNWrapper(opts); temp = fn(imdb, batch);
**function fn = getBatchSimpleNNWrapper(opts):**
    -> fn = @(imdb, batch) getBatchSimpleNN(imdb, batch, opts);
**function [im,labels] = getBatchSimpleNN(imdb, batch, opts):**
    -> im = imdb_get_batch_bcnn(images, opts, ...
        'prefetch', nargout == 0);
**function imo = imdb_get_batch_bcnn(images, varargin):**
    -> opts(i) = vl_argparse(opts(i), {varargin{1}(i),varargin{2:end}});
**function [opts, args] = vl_argparse(opts, args, varargin):**
    -> if nargout == 1
        opts = vl_argparse(opts, vertcat(params,values)) ;
**function [opts, args] = vl_argparse(opts, args, varargin):**
    -> if nargout == 1
        error('Unknown parameter ''%s'', param) ;
```

What the hell is cropSize?

After a few tracing during debugging, I find that `cropSize` is introduced by `net.meta.normalization` (`imdb_bcnn_train_dag.m`: line 71) to `bopts` which will be passed to function call `getImageStats(imdb, bopts)` and further be passed to function `imdb_get_batch_bcmn(images, varargin)` via `bopts -> opts -> varargin`, and finally be passed to function `vl_argparse(opts, args, varargin)`.

Inspect a few more lines around `opts = vl_argparse(opts, vertcat(params,values))` in `vl_argparse.m`:

```
% ./matconvnet/matlab/vl_argparse.m

if recursive && isstruct(args{ai})
    params = fieldnames(args{ai})' ;
    values = struct2cell(args{ai})' ;
    if nargout == 1
        opts = vl_argparse(opts, vertcat(params,values)) ;
    else
        [opts, rest] = vl_argparse(opts, reshape(vertcat(params,values), 1, [])) ;
        args{ai} = cell2struct(rest(2:2:end), rest(1:2:end), 2) ;
        keep(ai) = true ;
    end
    ai = ai + 1 ;
    continue ;
end
```

where `recursive` is set to `true`, and `args` is passed in as a pack `{varargin{1}(i),varargin{2:end}}`:

```
**function imo = imdb_get_batch_bcnn(images, varargin):**
```

```
-> opts(i) = vl_argparse(opts(i), {varargin{1}(i),varargin{2:end}});
**function [opts, args] = vl_argparse(opts, args, varargin):**
```

Look one more step back, the `varargin` in function `imdb_get_batch_bcnn()` is passed in via `opts,'prefetch', nargout == 0` :

```
**function [im,labels] = getBatchSimpleNN(imdb, batch, opts):**
-> im = imdb_get_batch_bcnn(images, opts, ...
    'prefetch', nargout == 0);
```

Bingo! `opts` is the first element of `args` in `vl_argparse` . Therefore, at the very beginning when `ai = 1` , `isstruct(args{ai})` will also be `true` . So the next sentence is going to proceed:

```
opts = vl_argparse(opts, vertcat(params,values)) ;
```

What is `opts` , `params` and `values` when program execution reach here?

```
opts =

    imageSize: [227 227]
        border: [0 0]
    keepAspect: 1
    numAugments: 1
transformation: 'none'
    averageImage: []
    rgbVariance: [0x3 single]
    interpolation: 'bilinear'
        numThreads: 1
        prefetch: 0
            scale: 1

params =

    'imageSize'    'keepAspect'    'averageImage'    'border'    'cropSize'    'interpolation'
    'numThreads'
```

```
values =

    [1x4 double]    [          1]          []    [1x2 double]    [1x2 double]    'bilinear'
    [          12]
```

where `params` is the `fieldnames` of `args{1}` : `params = fieldnames(args{ai})` and `values` packs their corresponding values: `values = struct2cell(args{ai})` .

So the function `vl_argparse(opts, args, varargin)` executes again with parameters the same `opts` and a new `args = vertcat(params,values)` .

Now none of the elements in `args` is a struct. The execution will then reach the code block, where the ERROR occurs:

```
% ./matconvnet/matlab/vl_argparse.m

param = args{ai} ;
value = args{ai+1} ;
p = find(strcmpi(param, optNames)) ;
if numel(p) ~= 1
    if nargout == 1
        error('Unkno parameter ''%s''', param) ;
    else
        keep([ai,ai+1]) = true ;
```



```

        ai = ai + 2 ;
        continue ;
    end
end
end

```

This code block is in a `while` loop. It finds the elements in `params` which is the first row of `args` one by one in `optNames` at each iteration of the loop. So what is the `optNames` ? It is a pack of all the fieldnames of `opts` . Have a comparison:

```

params =

    'imageSize'    'keepAspect'    'averageImage'    'border'    'cropSize'    'interpolation'    '
numThreads'

optNames =

    'imageSize'    'border'    'keepAspect'    'numAugments'    'transformation'    'averageImage'
'rgbVariance'    'interpolation'    'numThreads'    'prefetch'    'scale'

```

All elements but `cropSize` in `params` appear in `optNames` !!

That's why this error occurs:

```

Error: Unkown parameter 'cropSize'

```

So the problem is two fold:

```

`cropSize` should not appear in `net.meta.normalization`, or
`cropSize` should be made to appear in `opts` of `vl_argparse`.

```

Since `net.meta.normalization` seems to be the last thing I should touch, I would try to set `cropSize` to `opts` firstly. To do this, Let's have a check where the `opts` of `vl_argparse` comes from, which means we should look into the execution stack to find out who calls function `vl_argparse` .

It is `imo = imdb_get_batch_bcnn(images, varargin) :`

```

% /bcnn/imdb_get_batch_bcnn.m
function imo = imdb_get_batch_bcnn(images, varargin)
% imdb_get_batch_bcnn Load, preprocess, and pack images for BCNN evaluation
% For asymmetric model, the function preprocesses the images twice for two networks
% separately.

% OUTPUT
% imo: a cell array where each element is a cell array of images.
%     For symmetric bcnn model, numel(imo) will be 1 and imo{1} will be a
%     cell array of images
%     For asymmetric bcnn, numel(imo) will be 2. imo{1} is a cell array containing the preprocessed
images for network A
%     Similarly, imo{2} is a cell array containing the preprocessed images for network B

%
% Copyright (C) 2015 Tsung-Yu Lin, Aruni RoyChowdhury, Subhransu Maji.
% All rights reserved.
%
% This file is part of the BCNN and is made available under
% the terms of the BSD license (see the COPYING file).
%
% This file modified from CNN_IMAGENET_GET_BATCH of MatConvNet

for i=1:numel(varargin{1})

```

```

opts(i).imageSize = [227, 227] ;
opts(i).border = [0, 0] ;
opts(i).keepAspect = true;
opts(i).numAugments = 1 ;
opts(i).transformation = 'none' ;
opts(i).averageImage = [] ;
opts(i).rgbVariance = zeros(0,3,'single') ;
opts(i).interpolation = 'bilinear' ;
opts(i).numThreads = 1 ;
opts(i).prefetch = false;
opts(i).scale = 1;
opts(i) = vl_argparse(opts(i), {varargin{1}(i),varargin{2:end}});

```

Oh — My — God !!

...

Sir, the source code is poisonous !!

...

The code has built-in bugs!

Anyway just add a line `opts(i).cropSize = 227/256;` and give it another shot.

```

opts(i).scale = 1;
opts(i).cropSize = 227/256; % added by jcsu, 2016/11/16
opts(i) = vl_argparse(opts(i), {varargin{1}(i),varargin{2:end}});

```

OK, Let's remove output directory `checkgpu` and see what will happen! '

```
>> run_experiments_bcnn_train
```

Matlab standard outputs:

```

dataset: classes: 200 in use. These are:
  1: 001.Black_footed_Albatross (train:   30, test:   30 total:   60)
  2: 002.Laysan_Albatross (train:   30, test:   30 total:   60)
  3: 003.Sooty_Albatross (train:   30, test:   28 total:   58)
  4: 004.Groove_billed_Ani (train:   30, test:   30 total:   60)
  ...

```

```

199: 199.Winter_Wren (train:   30, test:   30 total:   60)
200: 200.Common_Yellowthroat (train:   30, test:   30 total:   60)
Inf:      **total** (train: 11788, test: 11788 total: 11788)
dataset: there are 11788 images (5994 trainval 5794 test)
Initialization: extracting bcnn feature of batch 1/185
Initialization: extracting bcnn feature of batch 2/185
...

```

```

Initialization: extracting bcnn feature of batch 184/185
Initialization: extracting bcnn feature of batch 185/185

```

layer	0	1	2
type	input	conv	softmx
name	n/a	classifier	loss
----- ----- ----- -----			
support	n/a	1	1
filt dim	n/a	262144	n/a
filt dilat	n/a	1	n/a
num filts	n/a	200	n/a

stride	n/a	1	1
pad	n/a	0	0
-----	-----	-----	-----
rf size	n/a	1	1
rf offset	n/a	1	1
rf stride	n/a	1	1
-----	-----	-----	-----
data size	NaNxNaN	NaNxNaN	NaNxNaN
data depth	NaN	200	1
data num	256	256	1
-----	-----	-----	-----
data mem	NaN	NaN	NaN
param mem	n/a	200MB	0B

parameter memory|200MB (5.2e+07 parameters)|
data memory| NaN (for batch size 256)|

Where is the `vl_nnbilinearpool` , `vl_nnsqrt` and `vl_nnl2norm` layers?

Well, since the network has been constructed without any error or warning and the `conv` and `softmaxloss` layers appear, I should trust the program. Notice that layers 1 and 2 are built-in layers, while the others are self-defined, maybe the reason is we do not provide it a way to output the related information.

```
cnn_train: resetting GPU
Clearing mex files
  CUDADevice with properties:

      Name: 'GeForce GTX TITAN X'
      Index: 1
      ComputeCapability: '5.2'
      SupportsDouble: 1
      DriverVersion: 8
      ToolkitVersion: 7
      MaxThreadsPerBlock: 1024
      MaxShmemPerBlock: 49152
      MaxThreadBlockSize: [1024 1024 64]
      MaxGridSize: [2.1475e+09 65535 65535]
      SIMDWidth: 32
      TotalMemory: 1.2800e+10
      AvailableMemory: 1.2623e+10
      MultiprocessorCount: 24
      ClockRateKHz: 1076000
      ComputeMode: 'Default'
      GPUOverlapsTransfers: 1
      KernelExecutionTimeout: 0
      CanMapHostMemory: 1
      DeviceSupported: 1
      DeviceSelected: 1

train: epoch 01:  1/ 24: 471.8 (471.8) Hz objective: 5.298 top1err: 1.000 top5err: 0.980
train: epoch 01:  2/ 24: 582.1 (759.6) Hz objective: 5.298 top1err: 0.998 top5err: 0.979
train: epoch 01:  3/ 24: 634.4 (773.4) Hz objective: 5.297 top1err: 0.993 top5err: 0.966
train: epoch 01:  4/ 24: 661.0 (756.2) Hz objective: 5.297 top1err: 0.995 top5err: 0.967
train: epoch 01:  5/ 24: 756.4 (757.6) Hz objective: 5.298 top1err: 0.995 top5err: 0.971
...
...
train: epoch 01: 23/ 24: 761.8 (776.9) Hz objective: 5.290 top1err: 0.990 top5err: 0.963
train: epoch 01: 24/ 24: 761.3 (731.7) Hz objective: 5.290 top1err: 0.990 top5err: 0.962
val: epoch 01:  1/ 23: 840.6 (840.6) Hz objective: 5.234 top1err: 0.875 top5err: 0.711
val: epoch 01:  2/ 23: 822.7 (805.5) Hz objective: 5.236 top1err: 0.889 top5err: 0.740
```

```

val: epoch 01: 3/ 23: 823.1 (824.0) Hz objective: 5.237 top1err: 0.897 top5err: 0.736
val: epoch 01: 4/ 23: 826.1 (835.1) Hz objective: 5.236 top1err: 0.898 top5err: 0.729
...
...
val: epoch 25: 20/ 23: 839.8 (845.5) Hz objective: 3.568 top1err: 0.473 top5err: 0.203
val: epoch 25: 21/ 23: 840.4 (852.4) Hz objective: 3.565 top1err: 0.470 top5err: 0.203
val: epoch 25: 22/ 23: 840.7 (848.2) Hz objective: 3.570 top1err: 0.472 top5err: 0.204
val: epoch 25: 23/ 23: 840.8 (841.4) Hz objective: 3.573 top1err: 0.475 top5err: 0.206
train: epoch 26: 1/ 24: 609.7 (609.7) Hz objective: 3.146 top1err: 0.207 top5err: 0.043
train: epoch 26: 2/ 24: 679.4 (767.2) Hz objective: 3.117 top1err: 0.203 top5err: 0.059
train: epoch 26: 3/ 24: 705.6 (764.5) Hz objective: 3.114 top1err: 0.203 top5err: 0.061
train: epoch 26: 4/ 24: 722.1 (776.6) Hz objective: 3.102 top1err: 0.195 top5err: 0.058
...
...
train: epoch 31: 20/ 24: 762.8 (785.9) Hz objective: 2.797 top1err: 0.185 top5err: 0.051
train: epoch 31: 21/ 24: 763.6 (780.9) Hz objective: 2.796 top1err: 0.184 top5err: 0.051
train: epoch 31: 22/ 24: 764.2 (775.9) Hz objective: 2.797 top1err: 0.184 top5err: 0.051
train: epoch 31: 23/ 24: 765.1 (786.0) Hz objective: 2.795 top1err: 0.182 top5err: 0.051
train: epoch 31: 24/ 24: 763.9 (703.1) Hz objective: 2.796 top1err: 0.183 top5err: 0.051
val: epoch 31: 1/ 23: 830.9 (830.9) Hz objective: 3.241 top1err: 0.453 top5err: 0.227
val: epoch 31: 2/ 23: 836.7 (842.6) Hz objective: 3.260 top1err: 0.438 top5err: 0.203
val: epoch 31: 3/ 23: 837.6 (839.3) Hz objective: 3.263 top1err: 0.436 top5err: 0.191
val: epoch 31: 4/ 23: 838.1 (839.8) Hz objective: 3.272 top1err: 0.438 top5err: 0.194
val: epoch 31: 5/ 23: 841.5 (848.4) Hz objective: 3.280 top1err: 0.436 top5err: 0.196
...
...
val: epoch 40: 19/ 23: 796.7 (796.4) Hz objective: 2.996 top1err: 0.435 top5err: 0.178
val: epoch 40: 20/ 23: 793.6 (740.2) Hz objective: 2.999 top1err: 0.433 top5err: 0.178
val: epoch 40: 21/ 23: 794.1 (804.5) Hz objective: 2.997 top1err: 0.432 top5err: 0.177
val: epoch 40: 22/ 23: 792.2 (752.8) Hz objective: 3.003 top1err: 0.435 top5err: 0.178
val: epoch 40: 23/ 23: 792.5 (803.9) Hz objective: 3.003 top1err: 0.436 top5err: 0.179
train: epoch 41: 1/ 24: 648.2 (648.2) Hz objective: 2.318 top1err: 0.141 top5err: 0.023
train: epoch 41: 2/ 24: 690.4 (738.5) Hz objective: 2.348 top1err: 0.133 top5err: 0.033
train: epoch 41: 3/ 24: 709.1 (749.8) Hz objective: 2.332 top1err: 0.141 top5err: 0.042
train: epoch 41: 4/ 24: 720.9 (758.6) Hz objective: 2.360 top1err: 0.153 top5err: 0.045
train: epoch 41: 5/ 24: 758.4 (757.5) Hz objective: 2.336 top1err: 0.142 top5err: 0.039
train: epoch 41: 6/ 24: 758.1 (756.9) Hz objective: 2.331 top1err: 0.143 top5err: 0.041
...
...

```

Good, It seems to work quite well!

Questions remain:

1. Has the `bcnn` network been correctly constructed?
2. The training outputs a final model, is it the final results that can reproduce the paper's experimental results?
3. How to fine-tune the hyper-parameters?
4. How to use the final model?
5. How to transplant it to Caffe framework?
6. Installation of `Caffe` ...
7. ...